

# Embedded Systems

Semester 1, 2025

*Dr Chris Lehnert*

Tarang Janawalkar

This work is licensed under a Creative Commons  
“Attribution-NonCommercial-ShareAlike 4.0 International” license.



# Contents

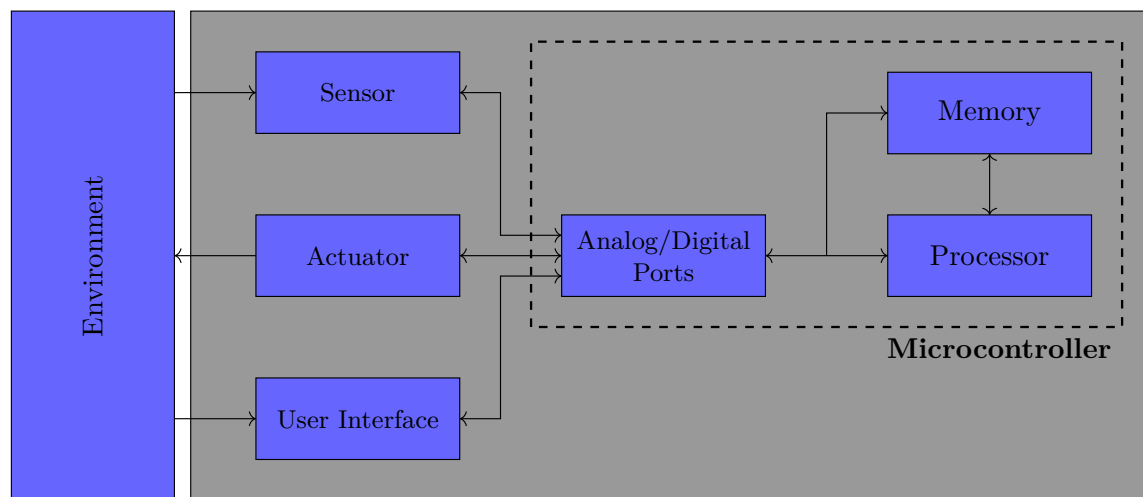
<b>Contents</b>	<b>1</b>
<b>1 Introduction</b>	<b>3</b>
1.1 Definition of an Embedded System . . . . .	3
1.1.1 Types of Embedded Systems . . . . .	3
1.2 Advanced RISC Machines . . . . .	3
1.2.1 ARM Cortex Cores . . . . .	4
1.3 Characteristics of an Embedded System . . . . .	4
1.3.1 Design Goals . . . . .	5
1.4 Real-Time Applications . . . . .	5
1.4.1 Real-Time Operating Systems . . . . .	5
1.5 Tiva C Series Microcontrollers . . . . .	5
1.5.1 Arm Cortex-M4 Processor . . . . .	6
1.5.2 Programming Models . . . . .	7
1.6 Microcontroller Architecture . . . . .	8
1.7 Microcontroller Memory . . . . .	8
1.7.1 Memory Map . . . . .	9
1.7.2 Bit-Banding . . . . .	9
<b>2 Microcontroller Peripherals</b>	<b>9</b>
2.1 System Control . . . . .	9
2.1.1 System Clock . . . . .	10
2.1.2 SysCtlClockFreqSet() . . . . .	11
2.1.3 Peripheral Enable . . . . .	11
2.2 General Purpose Input/Output (GPIO) . . . . .	12
2.2.1 GPIO Function API . . . . .	12
2.2.2 GPIO Peripheral Functions . . . . .	13
2.3 General Purpose Timers . . . . .	14
2.3.1 One-Shot & Continuous Mode . . . . .	14
2.3.2 Event Capture & Pulse Width Modulation . . . . .	14
2.3.3 General Purpose Timer Function API . . . . .	14
2.4 Watchdog Timer . . . . .	15
2.4.1 Interrupts . . . . .	15
2.4.2 Watchdog Timer Function API . . . . .	15
2.5 Hibernation . . . . .	15
2.6 Inter-Integrated Circuit (I2C) . . . . .	16
2.6.1 I2C Function API . . . . .	16
2.7 Analog to Digital Converter (ADC) . . . . .	16
2.7.1 ADC Function API . . . . .	17

<b>3</b>	<b>Exceptions</b>	<b>17</b>
3.1	Exceptions States . . . . .	17
3.2	Exception Types . . . . .	18
3.3	Exception Priorities . . . . .	18
3.4	Exception Processing . . . . .	19
3.5	Nested Vectored Interrupt Controller (NVIC) . . . . .	19
3.5.1	NVIC Triggering Mechanism . . . . .	20
3.5.2	Latency and Tail-Chaining . . . . .	20
<b>4</b>	<b>Multithreading</b>	<b>20</b>
4.1	Threads . . . . .	21
4.1.1	Thread Preemption . . . . .	21
4.1.2	Resource Sharing . . . . .	21
<b>5</b>	<b>FreeRTOS</b>	<b>22</b>
5.1	FreeRTOS Tasks . . . . .	22
5.1.1	Task Execution States . . . . .	23
5.2	FreeRTOS Scheduler . . . . .	23
5.2.1	Interrupts . . . . .	24
5.2.2	Tasks . . . . .	25
5.2.3	Idle Task . . . . .	25
5.3	FreeRTOS Tick . . . . .	25
5.4	FreeRTOS Kernel Configuration . . . . .	26
<b>6</b>	<b>Synchronisation</b>	<b>27</b>

# 1 Introduction

## 1.1 Definition of an Embedded System

An embedded system is a combination of computer hardware and software designed for a specific function or functions within a larger system. These systems typically contain computer hardware *within* their implementation and are used in devices to simplify system design and provide flexibility. Often, the user is unaware that a processor is present in the device as an embedded system comprises a suite of different components that communicate with each other to perform a specific task. These components include the processor, memory, and analog/digital ports—which form the microcontroller—and ports that are connected to various input/output devices such as sensors, actuators, and user interfaces—each of which interact with the environment. This is illustrated in the figure below, where the grey box represents the embedded system.



### 1.1.1 Types of Embedded Systems

Embedded systems can be classified into three main categories:

- **Centralised:** One node performs all work.
- **Distributed:** Nodes distribute work across sub-nodes.
- **Decentralised:** Nodes are only connected to peers in a network.

## 1.2 Advanced RISC Machines

Advanced RISC Machine (ARM) is a family of Instruction Set Architectures (ISAs) for computer processors. These ISAs are developed and designed by Arm Holdings so that they can be licensed to other companies that design their own ARM-based processors. ARM processors are found in many battery operated devices such as mobile phones, tablets, embedded systems, and some newer laptops.

Reduced Instruction Set Computer (RISC) processors are popular in such applications due to their high performance per watt and ability to execute all instructions in a single cycle. Additionally, because the architecture uses fixed-length instructions, instructions are also easier to pipeline, leading to increased parallelism. The RISC architecture focuses on small and highly-optimised instructions rather than the highly-specialised set of instructions found on Complex Instruction Set Computer (CISC) architectures such as x86. Although this may seem restrictive, this allows instructions to be executed at a greater frequency resulting in improved performance. Complex operations can then be performed in software using these instructions.

### 1.2.1 ARM Cortex Cores

ARM develops many families of processing cores for a range of different functions:

- **Cortex-A:** Highest performance cores—optimised for rich operating systems.
- **Cortex-R:** Fast response cores—optimised for high-performance, hard real-time applications.
- **Cortex-M:** High efficiency cores—optimised for discrete processing and microcontrollers.
- **SecurCore:** Tamper resistant cores—optimised for security applications.

## 1.3 Characteristics of an Embedded System

Embedded systems are characterised by several features. At a high level, they may be designed to be:

- Highly stable
- Time specific
- Task specific
- Cost effective
- Minimal in interface
- Easy to operate
- Real-time
- High-efficiency
- Reliable
- Memory constrained
- Power constrained
- Fault tolerant

### 1.3.1 Design Goals

These characteristics lead to several design goals in embedded systems such as:

- **Reliability:** Some systems may be critical to a mission, or life-threatening, and must be able to operate 24/7 without rebooting.
- **Performance:** Systems may need to respond to many events within a time frame using resources such as computing speed and power effectively. Constraints may need to be placed on inputs to prevent buffer overflows, and inaccuracies from floating-point calculations must be properly handled.
- **Cost:** Systems may be marketed to consumers and must therefore manufacturing minimise cost and be easy to produce.

## 1.4 Real-Time Applications

A system is said to be real-time if the total correctness of an operation not only depends on its logical correctness, but also upon the time in which it is performed. A primary design goal of real-time systems is **meeting deadlines**.

- **Soft real-time systems** execute as fast as possible requiring on explicit deadline on the response time.
- **Hard real-time systems** impose a strict deadline on the response time. If the deadline is missed, the system fails.

### 1.4.1 Real-Time Operating Systems

Embedded systems are typically developed using low-level programming languages such as C, C++, and assembly, for their performance and reliable compilation. The compilation process is different from that of a desktop application where code is compiled into an executable file which can be executed by the operating system. Instead, embedded systems (or those with sufficient resources) make use of **real-time kernel** libraries alongside application code to produce a single binary image that is flashed onto the device. These systems are known as real-time operating systems (RTOS). The kernel is software that manages this real-time system by providing abstractions for creating threads (tasks), scheduling, input/output operations, memory management, and other functions in an operating system.

## 1.5 Tiva C Series Microcontrollers

This unit uses the Texas Instruments Tiva C series TM4C1294NCPDT microcontroller which is housed on the EK-TM4C1294XL evaluation board. This microcontroller chip is based on an ARM Cortex-M4 core and includes several on-chip peripherals such as an Ethernet controller, USB interface, analog-to-digital converters (ADCs), and timers. The evaluation board also provides additional hardware such as LEDs, switches, a touch screen, and other input/output devices, all of which can be interfaced with the microcontroller.

### 1.5.1 Arm Cortex-M4 Processor

Some technical specifications of the Cortex-M4 processor are described below:

- **Architecture:** Armv7E-M
- **Bus Interface:** 3x Advanced Microcontroller Bus Architecture 3 Advanced High-performance Bus-Lite (AHBA 3 AHB-Lite) interface (Harvard bus architecture)
- **Instruction Set Architecture Support:** Thumb or Thumb-2<sup>1</sup>, hardware divide, single-cycle multiply, etc.
- **Pipeline:** 3-stage (Fetch-Decode-Execute)

A block diagram of the Cortex-M4 processor, and its register set are shown below.

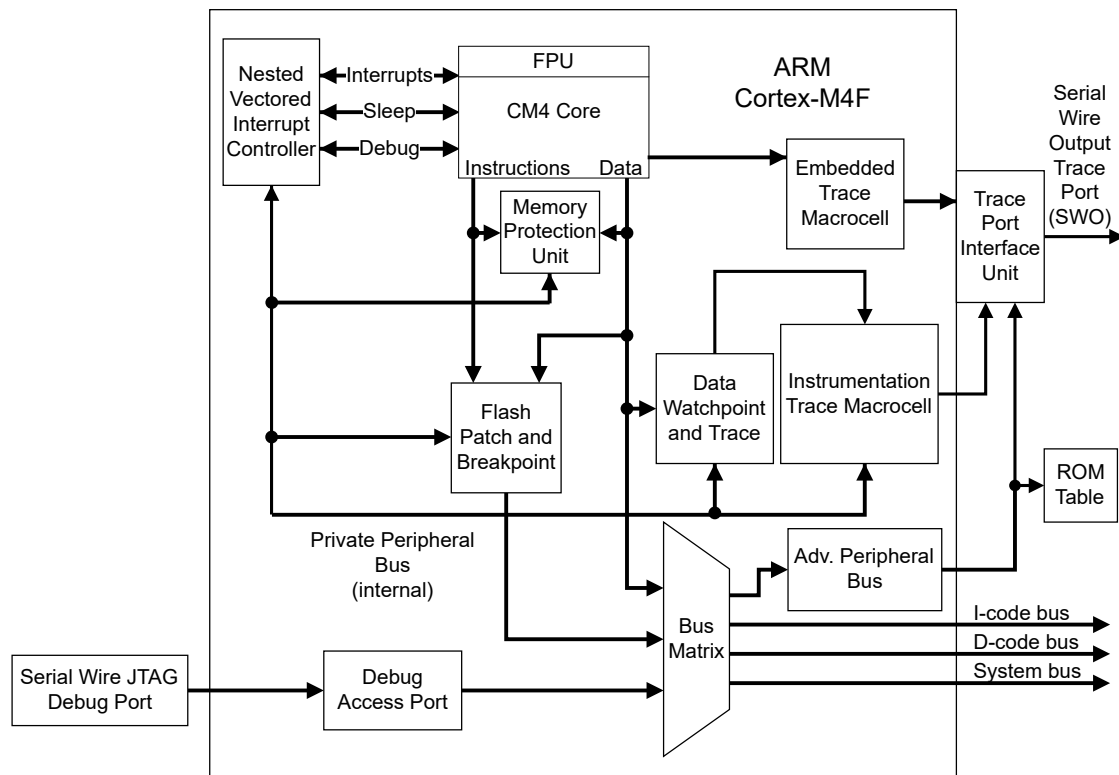


Figure 1: Block diagram of the Cortex-M4 processor.

<sup>1</sup>The thumb instruction set is a subset of the most commonly used 32-bit ARM instructions. Thumb instructions are 16 bits long and have corresponding 32-bit ARM instructions that perform the same function. They are used to reduce code size and improve performance in memory-constrained applications. Thumb2 provides enhancements to Thumb by introducing a new conditional instruction amongst other improvements.




Figure 2: Register set of the Cortex-M4 processor.

### 1.5.2 Programming Models

Tiva C series microcontrollers can be programmed using the **Direct Register Access Model** where the application accesses hardware registers directly using pointers and bitwise operations. This model results in small and more efficient code.

The registers in this model can be accessed by including the `tm4c1294ncpdt.h` header file, which contains register definitions for all peripherals on the Tiva C series microcontroller.

---

```
1 #include "inc/tm4c1294ncpdt.h"
```

---

The macros in this header file use the following naming conventions:

- **Suffixes:**

- `_R`: Access to the memory mapped register.



- `_M`: Mask for a multi-bit field.
- `_S`: Shift value for field alignment.
- **Register Name Structure**: `<MODULE><INSTANCE>_<REGISTER>_R`
- **Bit Field Name Structure**: `<MODULE><INSTANCE>_<REGISTER>_<FIELD>`. Bit fields with multiple values are often suffixed with `_M`, `_S`, or a value.

Alternatively, we can use the **Software Driver Model** where the application uses a development API such as the **TivaWare** software development kit (SDK) to access hardware registers. This model aims to simplify the development process by providing functions for accessing peripherals such as GPIO, UART, I2C, SPI, and ADC.

## 1.6 Microcontroller Architecture

A microcontroller is made up of a microprocessor, memory, peripherals, and I/O. Communication between the microprocessor and these devices is facilitated through the system bus that consists of an address bus, data bus, and control bus<sup>2</sup>. This bus allows shared communication between a single processor and device at a time. To overcome this limitation and allow multiple devices to communicate with multiple processors concurrently, a bus matrix is often used instead of, or in addition to, the system bus. Furthermore, microprocessor architecture also determines whether code memory and data memory are accessed via the same data bus, where:

- **von Neumann Bus Architecture** accesses code and data memory from the same bus.
- **Harvard Bus Architecture** accesses code and data memory from two separate buses, called the instruction code bus (I-code bus) and the data code bus (D-code bus).

## 1.7 Microcontroller Memory

The TM4C1294NCPDT microcontroller is integrated with the following set of on-chip memory:

- **Non-Volatile Memory** (retains data when powered off):
  - 1024 K B **Flash Memory** ( $4 \times 256$  K B banks) used for storing program code.
  - 6 K B **Electrically Erasable Programmable Read-Only Memory (EEPROM)** used for storing non-volatile data.
  - **Internal Read-Only Memory (ROM)** loaded with TivaWare for C Series software: TivaWare Peripheral Driver Library and TivaWare Boot Loader.
- **Volatile Memory** (loses data when powered off):
  - 256 K B **Single-Cycle Static Random Access Memory (SRAM)** used for very fast data access and frequent read/write operations. This memory is the runtime memory for the program stack, peripheral registers, and runtime variables.

---

<sup>2</sup>A bus refers to a group of lines carrying digital signals.

### 1.7.1 Memory Map

Desktop architectures typically have separate address spaces for memory and peripherals, where input and output devices are mapped to a separate address space. However, the TM4C1294NCPDT microcontroller uses a shared address space for memory and peripherals, which is known as the **memory map**.

Address Range	Memory Region	Description
0x0000_0000 - 0x1FFF_FFFF	Code	This executable region is for program code. Data can also be stored here.
0x2000_0000 - 0x3FFF_FFFF	SRAM	This executable region is for data. Code can also be stored here. This region includes bit band and bit band alias areas.
0x4000_0000 - 0x5FFF_FFFF	Peripheral	This region includes bit band and bit band alias areas.
0x6000_0000 - 0x9FFF_FFFF	External RAM	This executable region is for data.
0xA000_0000 - 0xDFFF_FFFF	External device	This region is for external device memory.
0xE000_0000 - 0xE00F_FFFF	Private peripheral bus	This region includes the NVIC, system timer, and system control block.
0xE010_0000 - 0xFFFF_FFFF	Reserved	-

Table 1: Memory Access Map on the TM4C1294NCPDT Microcontroller.

### 1.7.2 Bit-Banding

The ARM Cortex-M4 processor supports a feature called **bit-banding** which maps a full word of memory to a single bit in the bit-band region. This eliminates the need for read-modify-write operations, as individual bits can be set, cleared, or toggled directly from an alias address. This technique is used on the TM4C1294NCPDT as it uses a 32-bit word size, resulting in 4 G B of addressable memory, of which, peripherals only use a small portion. The remaining region of memory is therefore used for alias addresses that serve this purpose.

## 2 Microcontroller Peripherals

The following sections highlight the functionality of various peripherals and provide examples of their configuration using the TivaWare Peripheral Driver Library. All information in this section is taken from the User's Guide document.

### 2.1 System Control

System control determines the overall operation of the device by:

- controlling the system clock,
- configuring which peripherals are enabled,
- configuring the device and its resets, and by
- providing information about the device.

### 2.1.1 System Clock

The main system clock is used to clock the processor and all peripherals on the device. The system clock frequency is determined through the following steps:

1. **Select the input source:**

Choose an oscillator source (internal oscillator or external crystal) and configure the frequency of the oscillator.

2. **Apply a frequency multiplier:** (optional)

A Phase-Locked Loop (PLL) can be used to multiply an input frequency using a Voltage Controlled Oscillator (VCO).

3. **Apply Frequency Divider:** (optional)

The system clock divider can be used to divide the output frequency of the PLL or oscillator source.

4. **Select the system clock source:**

Choose whether the system clock is driven by the PLL output or directly by the oscillator source.

### Oscillator Source

The oscillator source can be one of the following:

- `SYSCTL_OSC_MAIN` to use an external crystal or oscillator.
- `SYSCTL_OSC_INT` to use the 16 MHz precision internal oscillator.
- `SYSCTL_OSC_INT30` to use the internal low frequency oscillator.
- `SYSCTL_OSC_EXT32` to use the hibernate modules 32.786 kHz oscillator.

When using an external crystal, the frequency is set using the macro `SYSCTL_XTAL_<frequency>MHZ` where `<frequency>` is the frequency of the crystal in MHz.

### System Clock Source

The system clock source may be configured to use the output of the PLL or be directly driven by the oscillator source using one of the following macros:

- `SYSCTL_USE_PLL` is used to select the PLL output as the system clock.
- `SYSCTL_USE_OSC` is used to choose one of the oscillators as the system clock.

When using the PLL, the VCO frequency can be configured using one of the following macros:

- `SYSCTL_CFG_VCO_480` to set the PLL VCO output to 480 MHz
- `SYSCTL_CFG_VCO_320` to set the PLL VCO output to 320 MHz

The `SysCtlClockFreqSet()` function attempts to match the requested system clock frequency to the closest possible value based on the configuration of the device clocking.

### 2.1.2 SysCtlClockFreqSet()

**Prototype:**

---

```
1 uint32_t SysCtlClockFreqSet(uint32_t ui32Config, uint32_t ui32SysClock);
```

---

**Parameters:**

- `ui32Config` is the required configuration of the device clocking.
- `ui32SysClock` is the requested processor frequency.

**Returns:**

The actual configured system clock frequency in Hz or zero if the value could not be changed due to a parameter error or PLL lock failure.

**Example:**

---

```
1 #include "inc/tm4c1294ncpdt.h"
2
3 int main(void) {
4     // Use a 25 MHz external crystal oscillator with a PLL VCO output of
5     // 480 MHz.
6     // The desired system clock frequency is 120 MHz which will require
7     // a system clock divider of 4.
8     SysCtlClockFreqSet(SYSCTL_XTAL_25MHZ | SYSCTL_OSC_MAIN |
9                         SYSCTL_USE_PLL | SYSCTL_CFG_VCO_480, 120000000);
10
11     // ...
12 }
```

---

### 2.1.3 Peripheral Enable

Most peripherals are disabled by default to reduce power consumption. A peripheral can be enabled using the system control module.

**SysCtlPeripheralEnable()**

**Prototype:**

---

```
1 void SysCtlPeripheralEnable(uint32_t ui32Peripheral)
```

---

**Parameters:**

- `ui32Peripheral` is the peripheral to enable.

**Returns:** None.

**Example:**

---

```
1  #include "inc/tm4c1294ncpdt.h"
2
3  int main(void) {
4      // Enable the GPIO Port A peripheral.
5      SysCtlPeripheralEnable(SYSCTL_PERIPH_GPIOA);
6
7      // ...
8  }
```

---

## 2.2 General Purpose Input/Output (GPIO)

The General Purpose Input/Output (GPIO) module provides control over GPIO ports and pins on the device. Each pin has the following capabilities:

- Can be configured as an input or an output. On reset, GPIOs default to being inputs.
- In input mode, can generate interrupts on high level, low level, rising edge, falling edge, or both edges.
- In output mode, can be configured for 2 mA, 4 mA, or 8 mA drive strength. The 8 mA drive strength configuration has optional slew rate control to limit the rise and fall times of the signal.

On reset, GPIOs default to 2 mA drive strength.

- Optional weak pull-up or pull-down resistors. On reset, GPIOs default to no pull-up or pulldown resistors.
- Optional open-drain operation. On reset, GPIOs default to standard push/pull operation.
- Can be configured to be a GPIO or a peripheral pin. On reset, the default is GPIO. Note that not all pins on all parts have peripheral functions, in which case the pin is only useful as a GPIO.

### 2.2.1 GPIO Function API

The GPIO API is broken into three groups of functions:

- The GPIO pins are configured with `GPIDirModeSet()`, `GPIOPadConfigSet()`, and `GPIOPinConfigure()`. The configuration can be read back with `GPIDirModeGet()` and `GPIOPadConfigGet()`.
- The GPIO pin state is accessed with `GPIOPinRead()` and `GPIOPinWrite()`.
- The GPIO interrupts are handled with `GPIOIntTypeSet()`, `GPIOIntTypeGet()`, `GPIOIntEnable()`, `GPIOIntDisable()`, `GPIOIntStatus()`, `GPIOIntClear()`, `GPIOIntRegister()`, and `GPIOIntUnregister()`.

### 2.2.2 GPIO Peripheral Functions

Many GPIO pins have other peripheral functions that can also use the GPIO pins for peripheral pins. Several convenience functions are provided to configure the pins in the required or recommended input/output configuration. These functions are:

- `GPIOPinTypeADC()`
- `GPIOPinTypeCAN()`
- `GPIOPinTypeComparator()`
- `GPIOPinTypeEPI()`
- `GPIOPinTypeEthernetLED()`
- `GPIOPinTypeEthernetMII()`
- `GPIOPinTypeGPIOInput()`
- `GPIOPinTypeGPIOOutput()`
- `GPIOPinTypeGPIOOutputOD()`
- `GPIOPinTypeI2C()`
- `GPIOPinTypeI2CSCL()`
- `GPIOPinTypeLCD()`
- `GPIOPinTypePWM()`
- `GPIOPinTypeQEI()`
- `GPIOPinTypeSSI()`
- `GPIOPinTypeTimer()`
- `GPIOPinTypeUART()`
- `GPIOPinTypeUSBAnalog()`
- `GPIOPinTypeUSBDigital()`
- `GPIOPinTypeWakeHigh()`
- `GPIOPinTypeWakeLow()`
- `GPIOPinWakeStatus()`
- `GPIONMATriggerEnable()`
- `GPIONMATriggerDisable()`
- `GPIOADCTriggerEnable()`
- `GPIOADCTriggerDisable()`

## 2.3 General Purpose Timers

The timer module provides two half-width timers/counters that can be configured to operate independently as timers or event counters or to operate as a combined full-width timer or Real Time Clock (RTC). Some timers provide 16-bit half-width timers and a 32-bit full-width timer, while others provide 32-bit half-width timers and a 64-bit full-width timer. Two half-width timers provided by a timer module are referred to as TimerA and TimerB, and the full-width timer is referred to as TimerA.

### 2.3.1 One-Shot & Continuous Mode

When configured as either a full-width or half-width timer, a timer can be set up to run as a one-shot timer or a continuous timer. If configured in one-shot mode, the timer ceases counting when it reaches zero when counting down or the load value when counting up. If configured in continuous mode, the timer counts to zero (counting down) or the load value (counting up), then reloads and continues counting. When configured as a full-width timer, the timer can also be configured to operate as an RTC. In this mode, the timer expects to be driven by a 32.768 KHz external clock, which is divided down to produce 1 second clock ticks.

### 2.3.2 Event Capture & Pulse Width Modulation

When in half-width mode, the timer can also be configured for event capture or as a Pulse Width Modulation (PWM) generator. When configured for event capture, the timer acts as a counter. It can be configured to either count the time between events or the events themselves. The type of event being counted can be configured as a positive edge, a negative edge, or both edges. When a timer is configured as a PWM generator, the input signal used to capture events becomes an output signal, and the timer drives an edge-aligned pulse onto that signal.

### 2.3.3 General Purpose Timer Function API

The timer API is broken into three groups of functions:

- Timer configuration is handled by `TimerConfigure()`, which performs the high level setup of the timer module; that is, it is used to set up full- or half-width modes, and to select between PWM, capture, and timer operations. Timer control is performed by `TimerEnable()`, `TimerDisable()`, `TimerControlLevel()`, `TimerControlTrigger()`, `TimerControlEvent()`, `TimerControlStall()`, `TimerRTCEnable()`, and `TimerRTCDisable()`.
- Timer content is managed with `TimerLoadSet()`, `TimerLoadGet()`, `TimerLoadSet64()`, `TimerLoadGet64()`, `TimerPrescaleSet()`, `TimerPrescaleGet()`, `TimerMatchSet()`, `TimerMatchGet()`, `TimerMatchSet64()`, `TimerMatchGet64()`, `TimerPrescaleMatchSet()`, `TimerPrescaleMatchGet()`, `TimerValueGet()`, `TimerValueGet64()`, and `TimerSynchronize()`.
- The interrupt handler for the Timer interrupt is managed with `TimerIntRegister()` and `TimerIntUnregister()`. The individual interrupt sources within the timer module are managed with `TimerIntEnable()`, `TimerIntDisable()`, `TimerIntStatus()`, and `TimerIntClear()`.

## 2.4 Watchdog Timer

A watchdog timer module's function is to prevent system hangs. The watchdog timer module consists of a 32-bit down counter, a programmable load register, interrupt generation logic, and a locking register. Once the watchdog timer has been configured, the lock register can be written to prevent the timer configuration from being inadvertently altered.

### 2.4.1 Interrupts

The watchdog timer can be configured to generate an interrupt to the processor after its first timeout, and to generate a reset signal after its second timeout. The watchdog timer module generates the first timeout signal when the 32-bit counter reaches the zero state after being enabled; enabling the counter also enables the watchdog timer interrupt. After the first timeout event, the 32-bit counter is reloaded with the value of the watchdog timer load register, and the timer resumes counting down from that value. If the timer counts down to its zero state again before the first timeout interrupt is cleared, and the reset signal has been enabled, the watchdog timer asserts its reset signal to the system. If the interrupt is cleared before the 32-bit counter reaches its second timeout, the 32-bit counter is loaded with the value in the load register, and counting resumes from that value. If the load register is written with a new value while the watchdog timer counter is counting, then the counter is loaded with the new value and continues counting.

### 2.4.2 Watchdog Timer Function API

The Watchdog Timer API is broken into two groups of functions:

- The Watchdog Timer interrupts are handled by the `WatchdogIntRegister()`, `WatchdogIntUnregister()`, `WatchdogIntEnable()`, `WatchdogIntClear()`, and `WatchdogIntStatus()` functions.
- Status and configuration functions for the Watchdog Timer module are `WatchdogEnable()`, `WatchdogRunning()`, `WatchdogLock()`, `WatchdogUnlock()`, `WatchdogLockState()`, `WatchdogReloadSet()`, `WatchdogReloadGet()`, `WatchdogValueGet()`, `WatchdogResetEnable()`, `WatchdogResetDisable()`, `WatchdogStallEnable()`, and `WatchdogStallDisable()`.

## 2.5 Hibernation

The Hibernation module allows the software application to remove power from the microcontroller, and then be powered on later based on specific time or when the external WAKE pin is asserted. The API provides functions to configure wake conditions, manage interrupts, read status, save and restore program state information, and request hibernation mode.

Some of the features of the Hibernation module are:

- 32-bit real time clock, with 15-bit subseconds counter on some devices
- Internal low frequency oscillator
- Calendar mode for the hibernation counter
- Tamper detection and response
- Trim register for fine tuning the RTC rate



- One RTC match registers for generating RTC events
- External WAKE pin to initiate a wake-up
- External RST pin and/or four GPIO port pins as alternate wake-up sources.
- Maintain GPIO state during hibernation.
- Low-battery detection
- 16 32-bit words of battery-backed memory
- Programmable interrupts for hibernation events

## 2.6 Inter-Integrated Circuit (I2C)

The I2C master and slave modules provide the ability to communicate to other IC devices over an I2C bus. The I2C bus is specified to support devices that can both transmit and receive (write and read) data. Also, devices on the I2C bus can be designated as either a master or a slave. The Tiva I2C modules support both sending and receiving data as either a master or a slave, and also support the simultaneous operation as both a master and a slave. Finally, the Tiva I2C modules can operate at the following speeds: Standard (100 kbps), Fast (400 kbps), Fast plus (1 Mbps) and High Speed (3.33 Mbps).

### 2.6.1 I2C Function API

The I2C API is broken into three groups of functions:

- The I2C pins are configured with `I2CMasterInitExpClk()`, `I2CMasterEnable()`, `I2CMasterDisable()`, `I2CSlaveInit()`, and `I2CSlaveEnable()`. The configuration can be read back with `I2CMasterErr()` and `I2CSlaveStatus()`.
- The I2C master and slave data is accessed with `I2CMasterSlaveAddrSet()`, `I2CMasterControl()`, `I2CMasterDataGet()`, `I2CMasterDataPut()`, `I2CSlaveDataGet()`, and `I2CSlaveDataPut()`.
- The I2C master and slave interrupts are handled with `I2CIntRegister()`, `I2CIntUnregister()`, `I2CMasterIntEnable()`, `I2CMasterIntDisable()`, `I2CMasterIntClear()`, `I2CMasterIntStatus()`, `I2CSlaveIntEnable()`, `I2CSlaveIntDisable()`, `I2CSlaveIntClear()`, and `I2CSlaveIntStatus()`.

## 2.7 Analog to Digital Converter (ADC)

The ADC module provides the ability to convert an analog voltage signal to a digital value. The ADC module can be configured to perform single-ended or differential conversions, and can be configured to perform conversions on a single channel or multiple channels.

### 2.7.1 ADC Function API

The ADC API is broken into three groups of functions:

- The ADC sample sequencers are configured with `ADCSequenceConfigure()` and `ADCSequenceStepConfigure()`. They are enabled and disabled with `ADCSequenceEnable()` and `ADCSequenceDisable()`. The captured data is obtained with `ADCSequenceDataGet()`. Sample sequencer FIFO overflow and underflow is managed with `ADCSequenceOverflow()`, `ADCSequenceOverflowClear()`, `ADCSequenceUnderflow()`, and `ADCSequenceUnderflowClear()`.
- Hardware oversampling of the ADC is controlled with `ADCHardwareOversampleConfigure()`. Software oversampling of the ADC is controlled with `ADCSoftwareOversampleConfigure()`, `ADCSoftwareOversampleStepConfigure()`, and `ADCSoftwareOversampleDataGet()`.
- The processor trigger is generated with `ADCProcessorTrigger()`.
- The interrupt handler for the ADC sample sequencer interrupts are managed with `ADCIntRegister()` and `ADCIntUnregister()`. The sample sequencer interrupt sources are managed with `ADCIntDisable()`, `ADCIntEnable()`, `ADCIntStatus()`, and `ADCIntClear()`.

## 3 Exceptions

Embedded systems often need to handle external events from touch screen inputs, sensors, actuators, keyboards, etc. These events are often asynchronous, or independent of the CPU, and may not necessarily be periodic or align with CPU clock cycles. Additionally, certain tasks that rely on precise timing, or tasks that must be performed concurrently, can be challenging to implement with a single CPU. These requirements are met and addressed through the use of **exceptions** and **interrupts**.

Exceptions are conditions used to halt the normal sequential flow of instructions in a program. They are prioritised and handled by the processor and the Nested Vectored Interrupt Controller (NVIC). When an exception occurs, the processor state is automatically stored on the stack and automatically restored from the stack at the end of the Interrupt Service Routine (ISR). The vector is fetched in parallel to the state saving, enabling efficient interrupt entry. The processor also supports tail-chaining, which enables back-to-back interrupts to be performed without the overhead of state saving and restoration.

### 3.1 Exceptions States

Each exception is in one of the following states:

- **Inactive.** The exception is not active and not pending.
- **Pending.** The exception is waiting to be serviced by the processor. An interrupt request from a peripheral or from software can change the state of the corresponding interrupt to pending.
- **Active.** An exception that is being serviced by the processor but has not completed.

**Note:** an exception handler can interrupt the execution of another exception handler. In this case, both exceptions are in the active state.

- **Active and Pending.** The exception is being serviced by the processor, and there is a pending exception from the same source.

### 3.2 Exception Types

Exceptions can occur due to various conditions from hardware or software, and can be categorised into several types:

- **Reset:** A special exception that occurs on power up or warm reset.
- **Non-Maskable Interrupt (NMI):** The highest priority exception, other than reset, that cannot be masked or preempted by any other exception. This exception is triggered using the NMI signal or triggered using the Interrupt Control and State (INTCTRL) register.
- **Faults:**
  - **Hard Fault:** An exception that occurs during exception processing, or because an exception cannot be managed by any other exception mechanism.
  - **Memory Management Fault:** An exception that occurs because of a memory protection related fault, including access violation and no match. Determined by the Memory Protection Unit (MPU).
  - **Bus Fault:** An exception that occurs because of a memory-related fault for an instruction or data memory transaction such as a prefetch fault or a memory access fault.
  - **Usage Fault:** An exception that occurs due to an instruction execution error, such as an undefined instruction, illegal unaligned access (half-word memory access), an invalid state on instruction execution, or division by zero.
- **Supervisor Call (SVCall):** An exception triggered by the SVC instruction, often used by operating systems for access to kernel functions and device drivers.
- **Debug Monitor** - An exception caused by the debug monitor when not halting.
- **PendSV** - Pendable System Service, used for system-level service requests, often for context switching in operating systems.
- **SysTick** - System Tick Timer, used for periodic interrupts, often for system timing in operating systems.
- **Interrupt (IRQ)** - Asynchronous exceptions signalled by peripherals or software requests, managed by the NVIC.

### 3.3 Exception Priorities

All exception types have an associated priority level that determines the order in which they are serviced by the processor. In the exception model of the ARM Cortex-M4 processor, exceptions with higher priorities are allowed to preempt the execution of lower priority exceptions. The vector number of an exception is used to locate where it is defined in the vector table, which is a list of pointers to the exception handlers. This table is typically located at the address 0. A table of all exception types is shown below.

Exception Type	Vector Number	Priority
–	0	–
Reset	1	–3 (highest)
Non-Maskable Interrupt (NMI)	2	–2
Hard Fault	3	–1
Memory Management Fault	4	programmable
Bus Fault	5	programmable
Usage Fault	6	programmable
–	7-10	–
SVCall	11	programmable
Debug Monitor	12	programmable
–	13	–
PendSV	14	programmable
SysTick	15	programmable
Interrupts	16-113	programmable

### 3.4 Exception Processing

When an exception is triggered, the processor performs the following steps:

1. **Stacking:** The contents of R0-R3, R12, Link Register (LR), Program Counter (PC), and Program Status Register (PSR) are pushed onto the stack. These registers need not be saved by the exception handler itself.
2. **Vector Fetching:** The exception is identified via the PSR and the processor fetches the vector number of the exception from the vector table.
3. **Load Link Register (LR):** The current mode of the CPU and stacking type are stored in the first 5 bits of the LR for return.
4. **Enable Handler Mode:** The processor switches to handler mode, which is a privileged mode that allows access to all system resources.
5. **Execute Exception Handler:** The processor executes the exception handler function associated with the vector number of the exception.

The mode of the processor is determined by the value of the Program Status Register (PSR), which holds the current exception type. Here, a value of 0 indicates Thread mode, and any non-zero value indicates Handler mode. Thread mode is the normal execution mode of the program and is always entered after a reset. Handler mode is an always-privileged mode that is entered when an exception occurs. It always returns to Thread mode when finished.

### 3.5 Nested Vectored Interrupt Controller (NVIC)

The Nested Vectored Interrupt Controller (NVIC) is a hardware component for event driven and real-time system design. It controls all exceptions and interrupts in the system. The NVIC provides the ability to individually enable and disable interrupts from specific devices, and establishes relative priorities among various interrupts. Vectored refers to the ability for the NVIC to determine the

address of the exception handler for each interrupt, which is stored in the vector table. This allows the processor to quickly fetch the address of the exception handler without needing to poll the device to decode which interrupt has occurred. Nested means that the NVIC can handle exceptions of varying priorities, allowing higher priority exceptions to preempt lower priority exceptions. This is particularly useful in real-time systems where certain tasks must be completed before others. For example, interrupts are assigned a programmable priority level from the highest priority (0) (default), to the lowest priority (7).

### 3.5.1 NVIC Triggering Mechanism

The NVIC supports both **level-sense** and **pulse detection** interrupts. Level-sense interrupts are triggered as long as the interrupt signal from a device is held at an active level, while pulse interrupts are triggered only when the interrupt signal transitions from inactive to active. These interrupts are latched by the NVIC, so that they enter a pending state before they can be serviced. This allows the NVIC to handle multiple interrupts of varying priorities without losing any interrupts. The interrupt enters the active state when the processor begins executing its interrupt handler.

### 3.5.2 Latency and Tail-Chaining

The NVIC is designed to minimise interrupt latency, which is the time taken by the processor to stack the current state of the program before executing the interrupt handler. This process and its inverse (unstacking) both require 12 clock cycles to complete. In the event where multiple interrupts are pending, the NVIC can perform **tail-chaining**, which allows the processor to handle multiple interrupts uninterrupted without unnecessarily unstacking and re-stacking the processor state. This can be done within 6 clock cycles, rather than 24, which is a significant reduction in latency.

## 4 Multithreading

As a system becomes more complex and needs to handle multiple asynchronous events, it becomes difficult to manage the flow of control sequentially. Consider the following three programming models:

- **Single Thread Polling:** The program runs in a single thread and polls devices sequentially, processing each device in the same loop. This approach often leads to non-real-time operation when processing times are long. These long processing times can also result in missed events if not polled frequently enough.
- **Single Thread and Processing within Interrupts:** The program runs in a single thread and processes devices within interrupts as they become available. While this approach reduces the latency of processing events, long processing times can still delay the processing of other events. Hence, it is not advisable to perform long processing tasks within an interrupt.
- **Multithreading and Signalling from Interrupts:** The program runs in multiple threads that can be scheduled independently. When an interrupt occurs, it uses a signalling mechanism to notify a thread that it needs to process the event. This approach decouples the hard real-time requirements of the interrupt from the processing requirements of the thread. A

scheduler can then be used to run these threads concurrently (on a single core), to ensure that progress is made on all threads, and that high-priority threads are serviced before lower-priority threads.

## 4.1 Threads

A thread is an independent unit of execution that has its own:

- **Program Scope:** Local variables, stack, and registers.
- **Initialisation:** Each thread can execute its own initialisation code when it is created.
- **Execution Pattern:** Each thread can run to completion or waits in a loop until it is signalled to run again.

Threads are used to achieve event-driven execution by waiting for events or specific data conditions before executing and provide access to real-time kernel functions for scheduling and synchronisation. Other benefits of threads include the ability to design and test each thread independently, and the reduced overall complexity of a program.

### 4.1.1 Thread Preemption

Thread preemption is the ability for a thread to be interrupted by another thread of higher priority, without requiring cooperation from the interrupted thread. This allows the system to respond to high priority events in a timely manner, while resuming the interrupted thread afterwards. When the processor needs to switch from thread A to thread B, it must perform a **context switch**, which involves saving the context of a thread A and restoring the context of thread B. This context contains the thread's stack pointer, program counter, registers, and flags, to allow the thread to resume execution correctly.

Threads may also be non-preemptive, or cooperative, where threads must voluntarily yield control explicitly. In this approach, other threads will be blocked until the current thread yields control. This can lead to simple and predictable execution, but may lead to low responsiveness if a high priority thread is blocked by a low priority thread that does not yield control, or in the worst case, starvation, if a thread never yields control. Despite this, cooperative threads may outperform preemptive threads as fewer interrupts require context switching, scheduling is simpler, and there are fewer synchronisation issues.

### 4.1.2 Resource Sharing

When two threads need to access the same resource, such as a global variable, careful synchronisation is required to ensure that the resource being accessed is not modified by another thread while it is being used. This often arises due to the non-atomicity of operations in the C programming language, where simple read-modify-write operations may translate to more than one Assembly instruction. This may lead to race conditions or data corruption if two threads attempt to read/modify the same resource at the same time. The section of code that may create this synchronisation problem is referred to as the *critical section*. Several synchronisation mechanisms are available to ensure that the critical section of code that accesses a shared resource is executed by only one thread at a time. These mechanisms include:

- **Mutexes:** A mutual exclusion lock that allows only one thread to access a resource at a time. Other threads attempting to access the resource will block until the mutex is released.
- **Semaphores:** A signalling mechanism that allows threads to signal each other when a resource is available or when an event has occurred. Semaphores can be used to control access to a resource or to synchronise the execution of threads.

## 5 FreeRTOS

FreeRTOS is a lightweight, microcontroller agnostic, real-time operating system (RTOS) kernel designed for embedded systems. It is an open-source implementation and provides a library of services that can be utilised in a scalable manner. These include: memory management on the stack and heap, debugging support with object naming and trace hooks, scheduling of tasks and interrupt service routines, and synchronisation primitives such as semaphores, mutexes, queues, and events. FreeRTOS provides support for preemptive, cooperative, or hybrid scheduling, which determines whether high-priority tasks can preempt lower-priority tasks. As the kernel is object-based, all APIs operate on self-contained objects such as tasks or semaphores, and changes to an object do not affect other objects.

### 5.1 FreeRTOS Tasks

FreeRTOS provides an abstraction for threads called **tasks**. This primitive enables the implementation of multithreading in a microcontroller environment. A task is a thread of execution that has its own stack and program counter, and is designed to run concurrently using signalling mechanisms. FreeRTOS also implements an idle task that does nothing when no other tasks are pending or active. It should be noted that ISRs are processed outside the FreeRTOS scheduler and must therefore handle synchronisation explicitly. FreeRTOS tasks are created using the following function:

---

```

1  #include "FreeRTOS.h"
2  #include "task.h"
3
4  BaseType_t xTaskCreate(TaskFunction_t pvTaskCode,
5                        const char * const pcName,
6                        unsigned short usStackDepth,
7                        void *pvParameters,
8                        UBaseType_t uxPriority,
9                        TaskHandle_t *pxCreatedTask);

```

---

where the task is specified by the function pointer `pvTaskCode`. A small example of a task is shown below:

---

```

1  void vTask(void *pvParameters) {
2      // Task initialisation
3
4      // Task execution loop (runs at least once)
5      while (1) {

```

---

```
6      // Perform task operations
7
8      // Optionally block or yield to allow other tasks to run
9      vTaskDelay(pdMS_TO_TICKS(100)); // Delay for 100 ms
10  }
11
12  // Task cleanup
13  }
14
15  int main(void) {
16      // Create the task
17      xTaskCreate(vTask, "MyTask", configMINIMAL_STACK_SIZE, NULL,
18                  tskIDLE_PRIORITY + 1, NULL);
19
20      // Start the scheduler
21      vTaskStartScheduler();
22
23      // The program should never reach here
24      while (1);
25  }
```

---

### 5.1.1 Task Execution States

Tasks can be in one of the following states:

- **Initialisation:** The task enters a Ready state, signalling its readiness to execute.
- **Ready:** The task is ready to run and waiting for the scheduler to allocate processor time.
- **Running:** The task is currently executing on the processor.
- **Blocked:** The task is waiting for a resource or an event to occur. It will not receive processor time until the resource or event becomes available. See `vTaskDelay()`, `xSemaphoreTake()`, or `xQueueReceive()`.
- **Suspended:** The task is suspended and will not be scheduled to run until it is resumed. This can be done by the kernel or by the task itself. See `vTaskSuspend()` and `vTaskResume()`.

On each tick, the scheduler selects a task in the Ready state to run using a scheduling algorithm, and gives the illusion of concurrent execution on a single-core processor through rapid task switching.

## 5.2 FreeRTOS Scheduler

The FreeRTOS scheduler determines which task executes at any moment and can suspend or resume tasks multiple times. It employs scheduling algorithms for task selection that ensure “fair” processor time allocation in multi-user systems, and adopts specific strategies for real-time/embedded systems. Tasks may be suspended by the kernel or choose self-suspension to sleep, wait for resources, or events, during which they are inactive and receive no processor time. If no other tasks are allocated processor time, the scheduler will run the idle task. Upon the creation of tasks, it is important to



call the `vTaskStartScheduler()` function to start the scheduler. This function does not return a value and therefore no code should be placed after it.

### 5.2.1 Interrupts

FreeRTOS relies on ISRs to handle hardware interrupts which are processed outside the FreeRTOS scheduler to meet hard deadlines. For this reason, FreeRTOS provides a set of ISR-safe APIs that are specifically designed to be called from within an ISR. These APIs are used to signal tasks, suspend tasks, and perform other operations that are safe to call from within an ISR. Some of the key ISR-safe APIs include:

---

```

1 xSemaphoreGiveFromISR(SemaphoreHandle_t xSemaphore,
2                       BaseType_t *pxHigherPriorityTaskWoken);
3
4 xQueueSendFromISR(QueueHandle_t xQueue,
5                  const void *pvItemToQueue,
6                  BaseType_t *pxHigherPriorityTaskWoken);
7
8 xEventGroupSetBitsFromISR(EventGroupHandle_t xEventGroup,
9                          const EventBits_t uxBitsToSet,
10                         BaseType_t *pxHigherPriorityTaskWoken);
11
12 xTaskNotifyFromISR(TaskHandle_t xTaskToNotify,
13                   uint32_t ulValue,
14                   eNotifyAction eAction,
15                   BaseType_t *pxHigherPriorityTaskWoken);

```

---

In addition to the respective kernel object (semaphore, queue, event group, or task), these functions also take an additional parameter called `pxHigherPriorityTaskWoken`. This parameter is used to unblock a higher priority task that is currently waiting for the event to occur. If the thread that was executing before the interrupt was of:

- **lower priority** than the thread that is unblocked, the scheduler will perform a context switch to the higher priority thread immediately after the ISR returns. This ensures that the system responds quickly to high-priority events.
- **higher priority** than the thread that is unblocked, the scheduler will resume processing the thread that was preempted and schedule the lower priority thread afterwards.

As mentioned previously, it is important to ensure that ISRs return quickly to avoid blocking time-critical tasks from being executed. Additionally, careful management of kernel objects is crucial when handling interrupts of varying priorities, as the NVIC may preempt interrupts. An example of an ISR that uses a semaphore to signal a task to wake up is shown below:

---

```

1 void vISR(void) {
2     BaseType_t xHigherPriorityTaskWoken = pdFALSE;
3

```

---

---

```

4      // Clear the interrupt flag to acknowledge the interrupt
5
6      // Perform the necessary operations in the ISR
7
8      // Signal a task to wake up
9      xSemaphoreGiveFromISR(xSemaphore, &xHigherPriorityTaskWoken);
10
11     // If a higher priority task was woken, request a context switch
12     portYIELD_FROM_ISR(xHigherPriorityTaskWoken);
13 }
14
15 void vTask(void *pvParameters) {
16     // Task initialisation
17
18     // Task execution loop
19     while (1) {
20         // Wait for the semaphore to be given by the ISR
21         if (xSemaphoreTake(xSemaphore, portMAX_DELAY) == pdTRUE) {
22             // Perform task operations after being woken by the ISR
23         }
24     }
25 }

```

---

### 5.2.2 Tasks

Tasks have a lower priority than ISRs and should be used to perform all intensive computation to keep ISRs short. They can wait (block) for resources to become available using signalling semaphores, queues, or event groups. As shown in the examples above tasks are simply function that run within a specific context (priority, stack, registers, etc.).

### 5.2.3 Idle Task

The `vStartTaskScheduler()` function calls the startup routine for each task and jumps into an idle task loop. This task may be used to check for stack overflows in other tasks and can be used for any non-time-critical processing. This task is run indefinitely until it is preempted by another task.

## 5.3 FreeRTOS Tick

Tasks may sleep or block for specific durations to enable effective scheduling. FreeRTOS uses an interrupt to increment a tick counter which keep track of the time since the scheduler started. On each tick, the kernel checks for tasks to unblock or wake, and to trigger context switches if a higher priority task is ready to run. We can convert to ticks using the `pdMS_TO_TICKS()` macro, which converts milliseconds to ticks based on the tick rate configured in FreeRTOS.

## 5.4 FreeRTOS Kernel Configuration

Several configuration options are available to customise the FreeRTOS kernel to suit the needs of the application. These options can be found in the `FreeRTOSConfig.h` file, which is included in the FreeRTOS kernel source code. Some of the key configuration options include:

- `configUSE_PREEMPTION`: Controls preemptive scheduling.
- `configCPU_CLOCK_HZ`: The CPU clock frequency in Hz.
- `configMINIMAL_STACK_SIZE`: The minimum stack size for tasks.
- `configMAX_PRIORITIES`: Sets the maximum number of task priority levels.
- `configUSE_16_BIT_TICKS`: Configures tick counter width.
- `configUSE_<FEATURE>`: Enables or disables kernel features such as semaphores, mutexes, event groups, and timers.
- `configIDLE_SHOULD_YIELD`: Adjusts scheduler behaviour and resource usage.

---

```

1  #define configUSE_TIMERS                1
2  #define configTIMER_TASK_PRIORITY      2
3  #define configTIMER_QUEUE_LENGTH      10
4  #define INCLUDE_xTimerPendFunctionCall 1
5  #define configTIMER_TASK_STACK_DEPTH  256
6  #define configUSE_EVENT_GROUPS         1
7  #define configUSE_PREEMPTION           1
8  #define configUSE_IDLE_HOOK            1
9  #define configUSE_TICK_HOOK            0
10 #define configCPU_CLOCK_HZ              ((unsigned long)120000000)
11 #define configTICK_RATE_HZ              ((portTickType)1000)
12 #define configMINIMAL_STACK_SIZE        ((unsigned short)200)
13 #define configTOTAL_HEAP_SIZE           ((size_t)(20240))
14 #define configMAX_TASK_NAME_LEN        (12)
15 #define configUSE_TRACE_FACILITY        1
16 #define configUSE_16_BIT_TICKS         0
17 #define configIDLE_SHOULD_YIELD         0
18 #define configUSE_CO_ROUTINES           0
19 #define configUSE_MUTEXES               1
20 #define configUSE_RECURSIVE_MUTEXES    1
21 #define configCHECK_FOR_STACK_OVERFLOW  2
22
23 #define configMAX_PRIORITIES             (16)
24 #define configMAX_CO_ROUTINE_PRIORITIES (2)
25 #define configQUEUE_REGISTRY_SIZE       10

```

---

## 6 Synchronisation

A semaphore is a synchronisation primitive that allows threads to safely share resources by controlling access to them. A semaphore